

Dissect: architecture and workflow

A research paper is parsed into **elements**, and the element is the unit of everything after that. A paragraph, a table, a figure, a display equation and a reference entry are all rows in one table, which is what lets the reader overlay, All Views, retrieval and a citation all read the same rows, and what lets a citation point at a bounding box on a page rather than at text that has lost its origin.

15,899

LINES OF PYTHON

41

API ROUTES

17

ELEMENT KINDS

19

MODEL PROVIDERS

11

EMBEDDING BACKENDS

1 REQUEST LIFECYCLE

Every request crosses the same four boundaries. In development the browser talks to the backend **directly**, because Next's dev rewrite buffers server sent events: the first token of an answer arrived at 0.0s direct and 19.0s through the proxy, which reads as a broken feature rather than a proxy. In production Vercel serves `/api` same origin, so the base is empty.

